

## BUSINESS UNDERSTANDING

An association of home owners have approached seeking to know what features in a house could be remodelled to increase the house prices and what features they should add as they list their houses in the market.

## DATA UNDERSTANDING

The data is collected from the King County House Sales dataset and the column description is as follows:

- `id` - Unique identifier for a house
- `date` - Date house was sold
- `price` - Sale price (prediction target)
- `bedrooms` - Number of bedrooms
- `bathrooms` - Number of bathrooms
- `sqft_living` - Square footage of living space in the home
- `sqft_lot` - Square footage of the lot
- `floors` - Number of floors (levels) in house
- `waterfront` - Whether the house is on a waterfront
  - Includes Duwamish, Elliott Bay, Puget Sound, Lake Union, Ship Canal, Lake Washington, Lake Sammamish, other lake, and river/slough waterfronts
- `view` - Quality of view from house
  - Includes views of Mt. Rainier, Olympics, Cascades, Territorial, Seattle Skyline, Puget Sound, Lake Washington, Lake Sammamish, small lake / river / creek, and other
- `condition` - How good the overall condition of the house is. Related to maintenance of house.
  - Rated 1-5 from poor to very good
- `grade` - Overall grade of the house. Related to the construction and design of the house.
  - Rated 1-13 from poor to excellent
- `sqft_above` - Square footage of house apart from basement
- `sqft_basement` - Square footage of the basement
- `yr_built` - Year when house was built
- `yr_renovated` - Year when house was renovated
- `zipcode` - ZIP Code used by the United States Postal Service
- `lat` - Latitude coordinate

- long - Longitude coordinate
- sqft\_living15 - The square footage of interior housing living space for the nearest 15 neighbors
- sqft\_lot15 - The square footage of the land lots of the nearest 15 neighbors

## Data Preprocessing

This involves checking the data, renaming some values, cleaning it and handling any outliers

Import all necessary libraries

```
In [352]: import pandas as pd
import seaborn as sns
import numpy as np
import math
import scipy.stats as stats
import statsmodels.api as sm
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import mean_absolute_error, mean_squared_error
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('ggplot')
```

Retrieve data from the dataset and preview the data

```
In [353]: houses = pd.read_csv("data/kc_house_data.csv")
houses.head()
```

```
Out[353]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN	NONE	...	7 Average	1180	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	NO	NONE	...	7 Average	2170	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	NO	NONE	...	6 Low Average	770	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	NO	NONE	...	7 Average	1050	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	NO	NONE	...	8 Good	1680	

5 rows × 21 columns



Get the information on the shape of data and columns of the data

```
In [354]: houses.shape
```

```
Out[354]: (21597, 21)
```

In [355]: `houses.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

```
In [356]: # Check for missing data  
houses.isna().sum()
```

```
Out[356]: id                0  
date                0  
price               0  
bedrooms            0  
bathrooms           0  
sqft_living          0  
sqft_lot             0  
floors              0  
waterfront          2376  
view                63  
condition            0  
grade               0  
sqft_above           0  
sqft_basement        0  
yr_built             0  
yr_renovated         3842  
zipcode              0  
lat                  0  
long                 0  
sqft_living15         0  
sqft_lot15           0  
dtype: int64
```

Fill the missing values with modes of their respective columns as they are very few missing values

```
In [357]: houses.waterfront.fillna(houses["waterfront"].mode().max(), inplace=True)  
houses.waterfront.value_counts()
```

```
Out[357]: NO      21451  
YES        146  
Name: waterfront, dtype: int64
```

```
In [358]: houses.view.fillna(houses["view"].mode().max(),inplace=True)
houses.view.value_counts()
```

```
Out[358]: NONE          19485
AVERAGE          957
GOOD             508
FAIR             330
EXCELLENT        317
Name: view, dtype: int64
```

```
In [359]: houses.yr_renovated.fillna(houses["yr_renovated"].mode().max(),inplace=True)
houses.yr_renovated.value_counts()
```

```
Out[359]: 0.0          20853
2014.0         73
2003.0         31
2013.0         31
2007.0         30
...
1946.0         1
1959.0         1
1971.0         1
1951.0         1
1954.0         1
Name: yr_renovated, Length: 70, dtype: int64
```

Check for missing values

```
In [360]: houses.isna().sum()
```

```
Out[360]: id            0  
date            0  
price           0  
bedrooms        0  
bathrooms        0  
sqft_living      0  
sqft_lot         0  
floors           0  
waterfront       0  
view             0  
condition        0  
grade            0  
sqft_above       0  
sqft_basement    0  
yr_built         0  
yr_renovated     0  
zipcode          0  
lat              0  
long             0  
sqft_living15    0  
sqft_lot15       0  
dtype: int64
```

Check individual columns for any irregular data

```
In [361]: for k,v in houses.items():
          print(f"For {k} the value counts are:\n {houses[k].value_counts()}")
```

```
For id the value counts are:
795000620      3
1825069031      2
2019200220      2
7129304540      2
1781500435      2
..
7812801125      1
4364700875      1
3021059276      1
880000205       1
1777500160      1
Name: id, Length: 21420, dtype: int64
For date the value counts are:
6/23/2014      142
6/25/2014      131
6/26/2014      131
7/8/2014       127
4/27/2015      126
```

```
In [362]: # Replace ? in sqft_basement with the mode
houses.sqft_basement.replace("?",0.0,inplace=True)
houses.sqft_basement.value_counts()
```

```
Out[362]: 0.0      12826
0.0      454
600.0     217
500.0     209
700.0     208
...
1275.0      1
1913.0      1
225.0       1
2570.0      1
1281.0      1
Name: sqft_basement, Length: 304, dtype: int64
```



```
In [363]: # Replace the strings in grade with the integers
houses["grade"] = houses.grade.apply(lambda x: x.split()[0])
```

```
In [364]: # Change the waterfront column to numbers on a scale No - 0 and Yes - 1
houses.waterfront.replace({"NO":0,"YES":1}, inplace=True)
houses.waterfront.value_counts()
```

```
Out[364]: 0    21451
          1      146
          Name: waterfront, dtype: int64
```

```
In [365]: # Change the scale quality in condition from 0 - 4
scale_cond = {
    "Poor": 0,
    "Fair": 1,
    "Average": 2,
    "Good": 3,
    "Very Good": 4
}
houses.condition.replace(scale_cond,inplace=True)
houses.condition.value_counts()
```

```
Out[365]: 2    14020
          3     5677
          4     1701
          1      170
          0        29
          Name: condition, dtype: int64
```

```
In [366]: # Change the scale quality in view from 0 - 4
scale_view = {
    "NONE": 0,
    "FAIR": 1,
    "AVERAGE": 2,
    "GOOD": 3,
    "EXCELLENT": 4
}
houses.view.replace(scale_view,inplace=True)
houses.view.value_counts()
```

```
Out[366]: 0    19485
          2     957
          3     508
          1     330
          4     317
          Name: view, dtype: int64
```

```
In [367]: houses.drop((houses[houses['grade'] == 13].index) | (houses[houses['grade'] == 3].index), inplace = True)
houses.grade.value_counts()
```

```
Out[367]: 7     8974
          8     6065
          9     2615
          6     2038
         10     1134
         11     399
          5     242
         12      89
          4      27
         13      13
          3         1
          Name: grade, dtype: int64
```

Change some of the data types of specific columns

```
In [368]: # Convert date column from object to datetime format
houses.date = pd.to_datetime(houses.date)
```

```
In [369]: # Convert sqft_basement to an integer format
houses["sqft_basement"] = houses.sqft_basement.astype(float)
```

```
In [370]: houses.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21597 non-null  int64
1   date                  21597 non-null  datetime64[ns]
2   price                 21597 non-null  float64
3   bedrooms              21597 non-null  int64
4   bathrooms             21597 non-null  float64
5   sqft_living           21597 non-null  int64
6   sqft_lot              21597 non-null  int64
7   floors                21597 non-null  float64
8   waterfront            21597 non-null  int64
9   view                  21597 non-null  int64
10  condition             21597 non-null  int64
11  grade                 21597 non-null  object
12  sqft_above            21597 non-null  int64
13  sqft_basement         21597 non-null  float64
14  yr_built              21597 non-null  int64
15  yr_renovated          21597 non-null  float64
16  zipcode               21597 non-null  int64
17  lat                   21597 non-null  float64
18  long                  21597 non-null  float64
19  sqft_living15         21597 non-null  int64
20  sqft_lot15            21597 non-null  int64
dtypes: datetime64[ns](1), float64(7), int64(12), object(1)
memory usage: 3.6+ MB
```

```
In [371]: # Save to csv
houses.to_csv("data/cleaned_data.csv",index=False)
```

## Feature Exploration

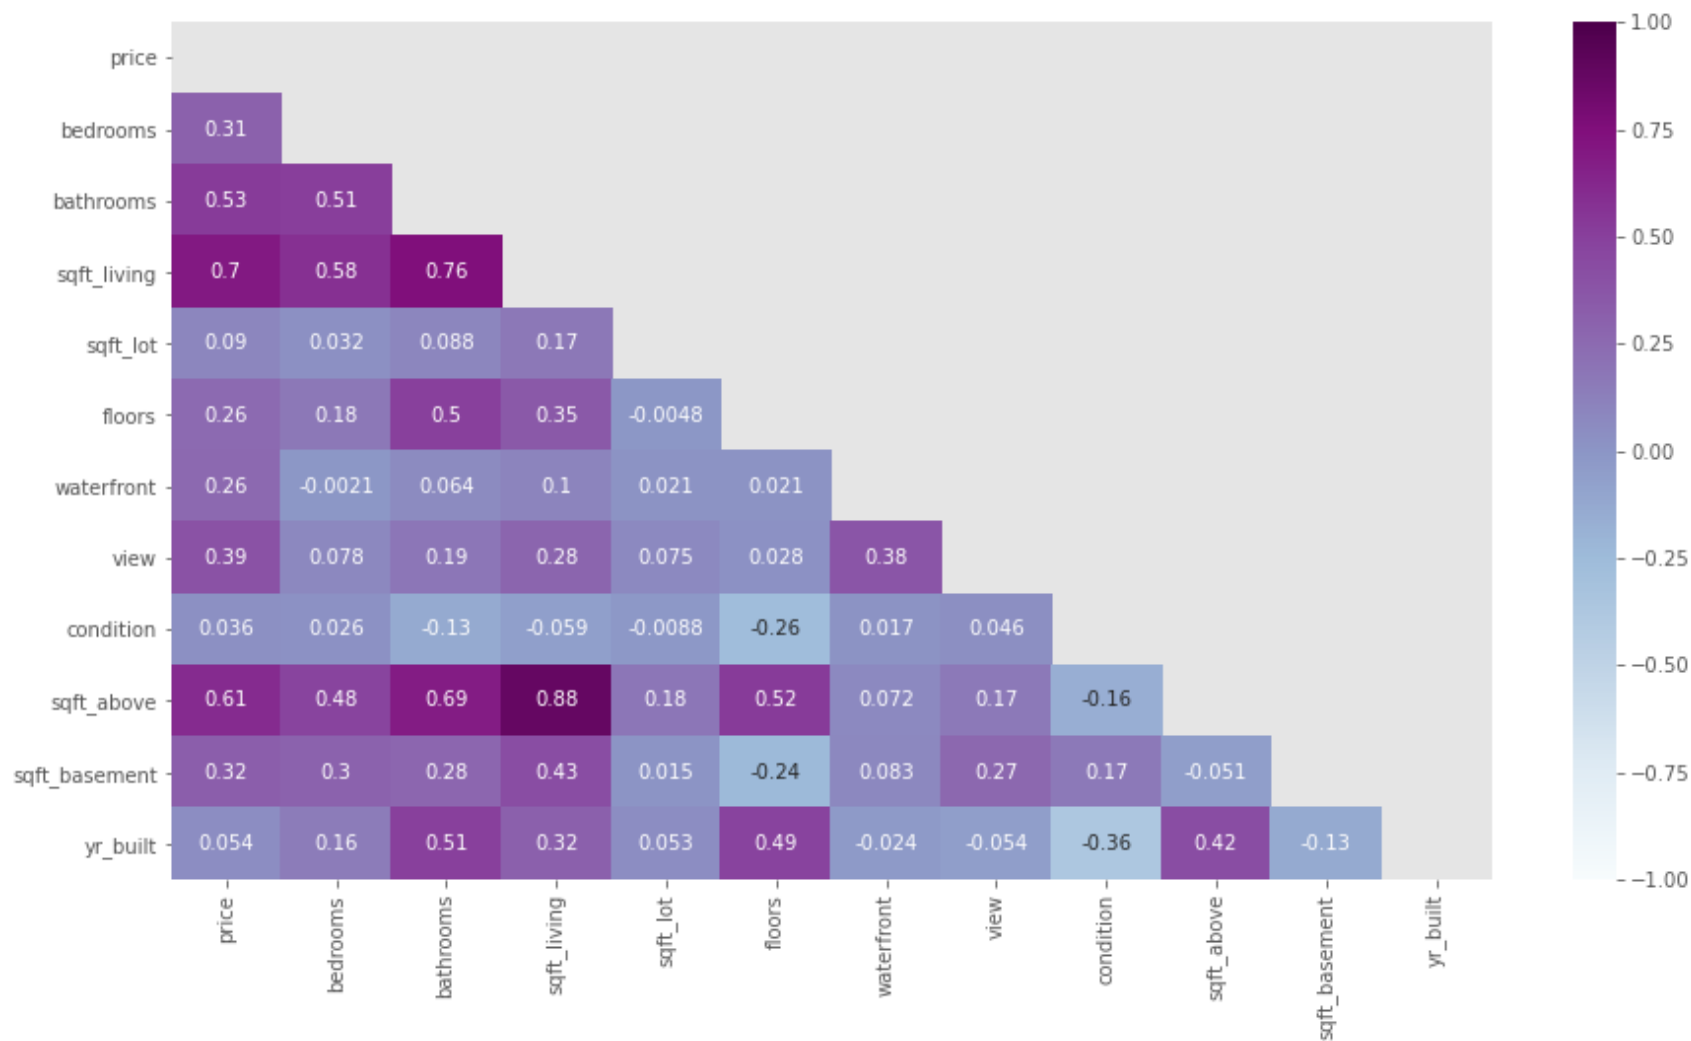
```
In [372]: # drop unwanted columns
houses.drop(["id", "zipcode", "long", "lat", "sqft_living15", "sqft_lot15", "yr_renovated", "date"],axis=1,inplace=True)
```

```
In [373]: houses.describe()
```

```
Out[373]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
<b>count</b>	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000	21597.000000	21597.000000	21597.000000
<b>mean</b>	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096	0.006760	0.233181	2.409825
<b>std</b>	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683	0.081944	0.764673	0.650546
<b>min</b>	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000	0.000000	0.000000	0.000000
<b>25%</b>	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000	0.000000	0.000000	2.000000
<b>50%</b>	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000	2.000000
<b>75%</b>	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000	0.000000	0.000000	3.000000
<b>max</b>	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000	4.000000	4.000000

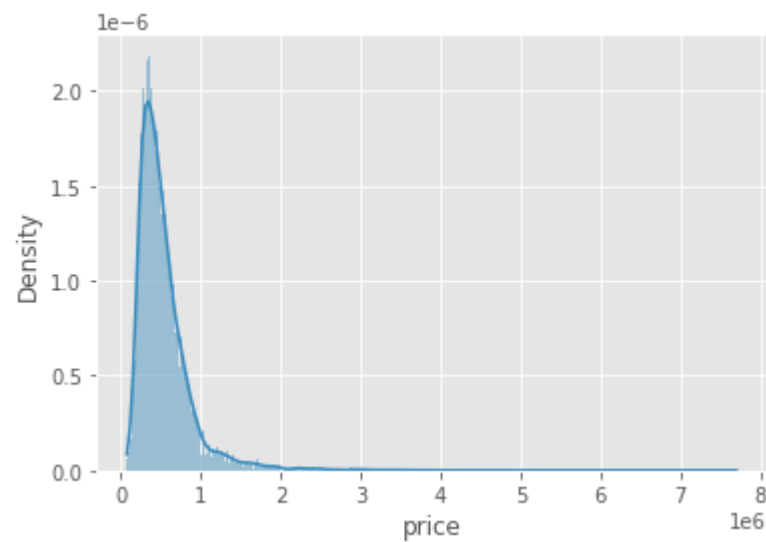
```
In [374]: # Display the correlation in a heatmap
fig,ax = plt.subplots(figsize=(15,8))
mask = np.triu(np.ones_like(houses.corr(), dtype=np.bool))
sns.heatmap(data=houses.corr(),center=0,vmin=-1,vmax=1,annot=True,mask=mask,cmap=sns.color_palette("BuPu", as_cmap=True))
plt.savefig("images/price heatmap correlation");
```



```
In [375]: # Correlation between other columns and price  
houses.corr()["price"]
```

```
Out[375]: price          1.000000  
bedrooms      0.308787  
bathrooms     0.525906  
sqft_living   0.701917  
sqft_lot      0.089876  
floors        0.256804  
waterfront    0.264306  
view          0.393497  
condition     0.036056  
sqft_above    0.605368  
sqft_basement 0.321108  
yr_built      0.053953  
Name: price, dtype: float64
```

```
In [376]: # Display a graph showing the price distribution
sns.histplot(data=houses, x="price", stat="density", kde=True)
plt.savefig("images/price density distribution");
```

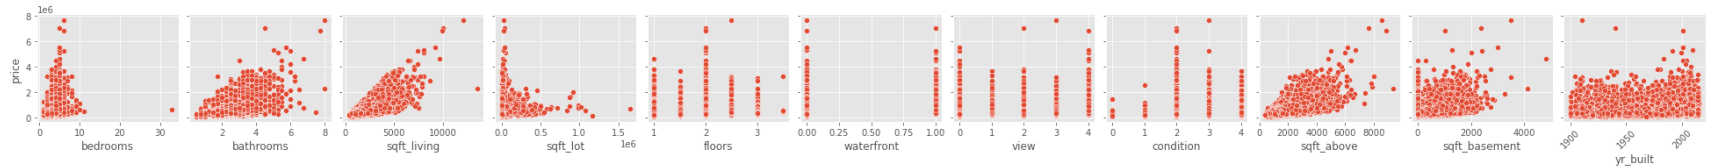


### Identifying variables

- Continuous numeric variables
- Discrete numeric variables
- String categorical variables
- Discrete categorical variables

```
In [377]: num_cols = houses.select_dtypes("number")
```

```
In [378]: sns.pairplot(data = houses,y_vars="price", x_vars=num_cols.columns.drop("price"),diag_kind = None)
plt.xticks(rotation=45)
plt.savefig("images/scatter plot distribution of numeric colums and price");
```



From the above plot we can see that the types of variables are:

- continuous numeric variables like sqft\_living,sqft\_lot,sqft\_above,sqft\_basement
- discrete numeric variables like bedrooms, bathrooms, floors, waterfront, view, condition, grade, yr\_built, yr\_renovated
- There are no string categoricals as all our values are of numerical type
- Discrete categorical variable like grade,floors,waterfront,view,condition

```
In [379]: houses.corr().abs()["price"]
```

```
Out[379]: price          1.000000
bedrooms        0.308787
bathrooms       0.525906
sqft_living     0.701917
sqft_lot        0.089876
floors          0.256804
waterfront      0.264306
view            0.393497
condition       0.036056
sqft_above      0.605368
sqft_basement   0.321108
yr_built        0.053953
Name: price, dtype: float64
```



```
In [380]: # Plot showing the category with the highest linearity with price
sns.scatterplot(x="sqft_living",y="price",data=houses)
plt.title("price vs sqft_living")
plt.savefig("images/price vs sqft_living");
```



## Modelling

```
In [381]: # Identify the selected features and create dummy variables
selected_features = houses.copy()

# Creating dummies
view_dummy = pd.get_dummies(selected_features, columns=["view"],drop_first=True)
grade_dummy = pd.get_dummies(selected_features, columns=["grade"],drop_first=True)
features_df = pd.concat([selected_features,grade_dummy,view_dummy],axis=1)
features_df = features_df.drop(['grade', 'view'], axis=1)
# Remove any duplicated columns
features_df = features_df.loc[:,~features_df.columns.duplicated()]

# Calculating the Mean Absolute Error of the model
def mae(x,y,model):
    y_pred = model.predict(sm.add_constant(x))
    mae = mean_absolute_error(y,y_pred)
    return mae
```

```
In [382]: # Check for correlation on features_df  
features_df.corr().abs()["price"]
```

```
Out[382]: price          1.000000  
bedrooms        0.308787  
bathrooms       0.525906  
sqft_living     0.701917  
sqft_lot        0.089876  
floors          0.256804  
waterfront     0.264306  
condition       0.036056  
sqft_above     0.605368  
sqft_basement   0.321108  
yr_built        0.053953  
grade_11        0.357589  
grade_12        0.291068  
grade_13        0.211806  
grade_3         0.005155  
grade_4         0.031618  
grade_5         0.084549  
grade_6         0.209463  
grade_7         0.316053  
grade_8         0.004576  
grade_9         0.235859  
view_1          0.092597  
view_2          0.147179  
view_3          0.182932  
view_4          0.303059  
Name: price, dtype: float64
```

### Baseline model

- sqft\_living is selected as the independent column and we want to use as it has the highest correlation with the price

```
In [383]: # Identify X and y variables that we shall use. sqft_living is selected as it has the highest  
X_model_baseline = features_df['sqft_living']  
y_model_baseline = features_df['price']  
  
# Create and fit the model  
model_baseline = sm.OLS(y_model_baseline, sm.add_constant(X_model_baseline)).fit()
```

```
In [384]: # Baseline model
model_baseline.summary()
```

Out[384]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.493
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.493
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	2.097e+04
<b>Date:</b>	Sat, 01 Oct 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	03:50:53	<b>Log-Likelihood:</b>	-3.0006e+05
<b>No. Observations:</b>	21597	<b>AIC:</b>	6.001e+05
<b>Df Residuals:</b>	21595	<b>BIC:</b>	6.001e+05
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-4.399e+04	4410.023	-9.975	0.000	-5.26e+04	-3.53e+04
<b>sqft_living</b>	280.8630	1.939	144.819	0.000	277.062	284.664

<b>Omnibus:</b>	14801.942	<b>Durbin-Watson:</b>	1.982
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	542662.604
<b>Skew:</b>	2.820	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	26.901	<b>Cond. No.</b>	5.63e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [385]: mae(X_model_baseline,y_model_baseline,model_baseline)
```

```
Out[385]: 173824.8874961748
```

The baseline model is interpreted as follows:

- The formula for getting price is written as

$$\text{price} = -43990 + 280 * \text{sqft\_living}$$

- The overall model explains about 49.3% of the variance in the prices.
- The sqft\_living coefficient is statistically significant.
- The model is off by about 173,824 in price \* The house price is 43,990 when the sqft\_living is 0
- For a one square foot increase in sqft\_living, there is a \$280 increase in price

### Model 1

- We shall perform multiple linear regression
- This will be a model that has 2 independent continuous numerical values

```
In [386]: X_model1 = features_df.loc[:,["sqft_living", "sqft_basement"]]  
y_model1 = features_df["price"]  
  
model1 = sm.OLS(y_model1, sm.add_constant(X_model1)).fit()
```

```
In [387]: # Model 1
model1.summary()
```

Out[387]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.493
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.493
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1.051e+04
<b>Date:</b>	Sat, 01 Oct 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	03:50:53	<b>Log-Likelihood:</b>	-3.0005e+05
<b>No. Observations:</b>	21597	<b>AIC:</b>	6.001e+05
<b>Df Residuals:</b>	21594	<b>BIC:</b>	6.001e+05
<b>Df Model:</b>	2		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-4.106e+04	4453.261	-9.220	0.000	-4.98e+04	-3.23e+04
<b>sqft_living</b>	276.6134	2.146	128.920	0.000	272.408	280.819
<b>sqft_basement</b>	20.6946	4.479	4.620	0.000	11.916	29.474

<b>Omnibus:</b>	14754.603	<b>Durbin-Watson:</b>	1.982
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	538977.524
<b>Skew:</b>	2.807	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	26.821	<b>Cond. No.</b>	5.75e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.75e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [388]: mae(X_model1, y_model1, model1)
```

```
Out[388]: 173566.09992442088
```

The first model is interpreted as follows:

- The formula for getting price is  $\text{price} = -41060 + 276.61(\text{sqft\_living}) + 20.70(\text{sqft\_basement})$
- The overall model explains about 49.3% of the variance in the prices.
- The model is off by about 173,566 in price \* *All the coefficients are statistically significant* \* *The house price is -41060 when the sqft\_living and sqft\_basement are 0*
- For a 1 square foot increase in the sqft\_living, there is a 276.61 increase in price \* *For a 1 square foot increase in the sqft\_basement, there is a 20.70 increase in price*

## Model 2

- Perform multiple linear regression on two other continuous variables

```
In [389]: X_model2 = features_df.loc[:, ["sqft_lot", "sqft_above"]]  
y_model2 = features_df["price"]  
  
model2 = sm.OLS(y_model2, sm.add_constant(X_model2)).fit()
```



```
In [390]: # Model 2
model2.summary()
```

Out[390]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.367
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.367
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	6259.
<b>Date:</b>	Sat, 01 Oct 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	03:50:54	<b>Log-Likelihood:</b>	-3.0245e+05
<b>No. Observations:</b>	21597	<b>AIC:</b>	6.049e+05
<b>Df Residuals:</b>	21594	<b>BIC:</b>	6.049e+05
<b>Df Model:</b>	2		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	5.948e+04	4736.364	12.559	0.000	5.02e+04	6.88e+04
<b>sqft_lot</b>	-0.1983	0.049	-4.058	0.000	-0.294	-0.103
<b>sqft_above</b>	270.4952	2.445	110.642	0.000	265.703	275.287

<b>Omnibus:</b>	16444.632	<b>Durbin-Watson:</b>	1.987
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	722025.818
<b>Skew:</b>	3.252	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	30.569	<b>Cond. No.</b>	1.05e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.05e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [391]: mae(X_model12,y_model12,model12)
```

```
Out[391]: 192035.56368124706
```

The second model is interpreted as follows:

- The formula for getting price is  $\text{price} = 276700 + 1161000(\text{waterfront}) + 171100(\text{floors})$
- The overall model explains about 0.133% of the variance in the prices.
- The model is off by about 220,007 *in price* \* *All the coefficients are statistically significant* \* *The house price is 276700* when the waterfront and floors are 0
- For an availability of a waterfront, there is a 1,161,000 *increase in price* \* *For a 1 increase in the floors, there is a 171,100 increase in price*

### Model 3

- The model will drop all the features that has the smallest correlation

```
In [392]: X_model3 = features_df.drop(["price"],axis=1)
X_model3.drop(["grade_3","grade_8"],axis=1,inplace=True)
y_model3 = features_df["price"]

model3 = sm.OLS(y_model3,sm.add_constant(X_model3)).fit()
```

```
In [393]: # Model 3
model3.summary()
```

Out[393]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.654
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.654
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1856.
<b>Date:</b>	Sat, 01 Oct 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	03:50:54	<b>Log-Likelihood:</b>	-2.9592e+05
<b>No. Observations:</b>	21597	<b>AIC:</b>	5.919e+05
<b>Df Residuals:</b>	21574	<b>BIC:</b>	5.921e+05
<b>Df Model:</b>	22		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	6.412e+06	1.35e+05	47.515	0.000	6.15e+06	6.68e+06
<b>bedrooms</b>	-3.738e+04	2053.371	-18.203	0.000	-4.14e+04	-3.34e+04
<b>bathrooms</b>	4.954e+04	3491.594	14.189	0.000	4.27e+04	5.64e+04
<b>sqft_living</b>	145.8608	19.383	7.525	0.000	107.869	183.852
<b>sqft_lot</b>	-0.2692	0.037	-7.343	0.000	-0.341	-0.197
<b>floors</b>	3.944e+04	3783.650	10.423	0.000	3.2e+04	4.69e+04
<b>waterfront</b>	5.178e+05	2.19e+04	23.639	0.000	4.75e+05	5.61e+05
<b>condition</b>	2.033e+04	2468.439	8.236	0.000	1.55e+04	2.52e+04
<b>sqft_above</b>	47.6731	19.355	2.463	0.014	9.736	85.610
<b>sqft_basement</b>	40.9728	19.235	2.130	0.033	3.271	78.674
<b>yr_built</b>	-3211.9697	68.223	-47.080	0.000	-3345.692	-3078.247
<b>grade_11</b>	4.301e+05	1.23e+04	35.104	0.000	4.06e+05	4.54e+05
<b>grade_12</b>	8.362e+05	2.44e+04	34.223	0.000	7.88e+05	8.84e+05

<b>grade_13</b>	1.923e+06	6.14e+04	31.307	0.000	1.8e+06	2.04e+06
<b>grade_4</b>	-1.893e+05	4.2e+04	-4.507	0.000	-2.72e+05	-1.07e+05
<b>grade_5</b>	-2.059e+05	1.48e+04	-13.914	0.000	-2.35e+05	-1.77e+05
<b>grade_6</b>	-1.607e+05	6609.464	-24.318	0.000	-1.74e+05	-1.48e+05
<b>grade_7</b>	-1.005e+05	4052.420	-24.795	0.000	-1.08e+05	-9.25e+04
<b>grade_9</b>	6.549e+04	5106.767	12.824	0.000	5.55e+04	7.55e+04
<b>view_1</b>	1.225e+05	1.21e+04	10.108	0.000	9.88e+04	1.46e+05
<b>view_2</b>	6.27e+04	7345.826	8.536	0.000	4.83e+04	7.71e+04
<b>view_3</b>	1.321e+05	1e+04	13.173	0.000	1.12e+05	1.52e+05
<b>view_4</b>	2.865e+05	1.52e+04	18.825	0.000	2.57e+05	3.16e+05

<b>Omnibus:</b>	11446.125	<b>Durbin-Watson:</b>	1.977
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	302660.740
<b>Skew:</b>	2.017	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	20.890	<b>Cond. No.</b>	4.05e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.05e+06. This might indicate that there are strong multicollinearity or other numerical problems.

In [394]: `mae(X_model3,y_model3,model3)`

Out[394]: 141119.16094729994

The third model is interpreted as follows:

- The overall model explains about 65.4% of the variance in the prices.
- The model is off by about \$141,119 in price
- All coefficients are statistically significant
- Compared to grade 3, grades 4,5,6 and 7 have a price decrease

- Compared to grade 3, grades 9,11,12 and 13 have a price increase with grade 13 having the most increase
- Compared to view 0, views 1,3 and 4 have a price increase with view 4 having the most

#### Model 4

Check for pairs that are highly correlated and remove some of the features

```
In [395]: # Checking for independent variables that are highly correlated and dropping them
def correlation(df, threshold):
    """
    The function takes in a dataframe and threshold
    The threshold determines what minimum correlation value you want to check from the dataframe
    It returns the columns that fit the threshold
    """

    corr_cols = set()
    corr_matrix = df.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i,j]) > threshold:
                new_col = corr_matrix.columns[i]
                corr_cols.add(new_col)
    return corr_cols
```

```
In [396]: X_model4 = features_df.drop(["price", "grade_3", "grade_8"], axis=1)
X_model4 = X_model4.drop(correlation(features_df, 0.6), axis=1)
y_model4 = features_df["price"]

model4 = sm.OLS(y_model4, sm.add_constant(X_model4)).fit()
```

```
In [397]: # Model 4
model14.summary()
```

Out[397]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.600
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.600
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1620.
<b>Date:</b>	Sat, 01 Oct 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	03:50:54	<b>Log-Likelihood:</b>	-2.9749e+05
<b>No. Observations:</b>	21597	<b>AIC:</b>	5.950e+05
<b>Df Residuals:</b>	21576	<b>BIC:</b>	5.952e+05
<b>Df Model:</b>	20		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	6.756e+06	1.45e+05	46.641	0.000	6.47e+06	7.04e+06
<b>bedrooms</b>	5441.2002	2060.320	2.641	0.008	1402.820	9479.580
<b>bathrooms</b>	1.254e+05	3474.270	36.089	0.000	1.19e+05	1.32e+05
<b>sqft_lot</b>	0.1213	0.039	3.131	0.002	0.045	0.197
<b>floors</b>	6.153e+04	4031.065	15.264	0.000	5.36e+04	6.94e+04
<b>waterfront</b>	5.488e+05	2.35e+04	23.307	0.000	5.03e+05	5.95e+05
<b>condition</b>	1.866e+04	2653.117	7.032	0.000	1.35e+04	2.39e+04
<b>sqft_basement</b>	81.6472	4.494	18.170	0.000	72.840	90.455
<b>yr_built</b>	-3358.4918	73.246	-45.852	0.000	-3502.059	-3214.924
<b>grade_11</b>	6.722e+05	1.24e+04	54.262	0.000	6.48e+05	6.96e+05
<b>grade_12</b>	1.204e+06	2.54e+04	47.463	0.000	1.15e+06	1.25e+06
<b>grade_13</b>	2.531e+06	6.51e+04	38.874	0.000	2.4e+06	2.66e+06

<b>grade_4</b>	-2.829e+05	4.51e+04	-6.269	0.000	-3.71e+05	-1.94e+05
<b>grade_5</b>	-2.971e+05	1.58e+04	-18.781	0.000	-3.28e+05	-2.66e+05
<b>grade_6</b>	-2.357e+05	6969.070	-33.826	0.000	-2.49e+05	-2.22e+05
<b>grade_7</b>	-1.627e+05	4202.279	-38.719	0.000	-1.71e+05	-1.54e+05
<b>grade_9</b>	1.276e+05	5368.923	23.768	0.000	1.17e+05	1.38e+05
<b>view_1</b>	1.42e+05	1.3e+04	10.899	0.000	1.16e+05	1.68e+05
<b>view_2</b>	7.985e+04	7886.736	10.124	0.000	6.44e+04	9.53e+04
<b>view_3</b>	1.56e+05	1.08e+04	14.477	0.000	1.35e+05	1.77e+05
<b>view_4</b>	3.25e+05	1.63e+04	19.879	0.000	2.93e+05	3.57e+05

<b>Omnibus:</b>	13378.851	<b>Durbin-Watson:</b>	1.969
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	445620.882
<b>Skew:</b>	2.450	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	24.707	<b>Cond. No.</b>	4.04e+06

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.04e+06. This might indicate that there are strong multicollinearity or other numerical problems.

In [398]: `mae(X_model14,y_model14,model14)`

Out[398]: 149953.7152730511

The fourth model is interpreted as follows:

- The overall model explains about 60% of the variance in the prices.
- The model is off by about \$149,953 in price
- All coefficients are statistically significant
- Compared to grade 3, grades 4,5,6 and 7 have a price decrease

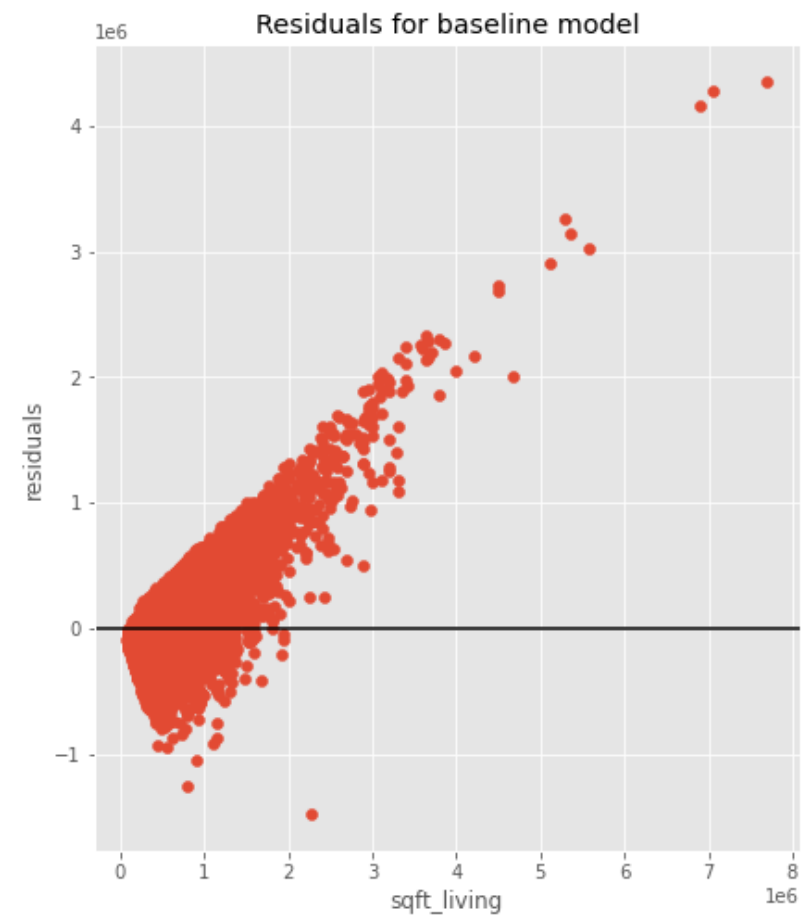
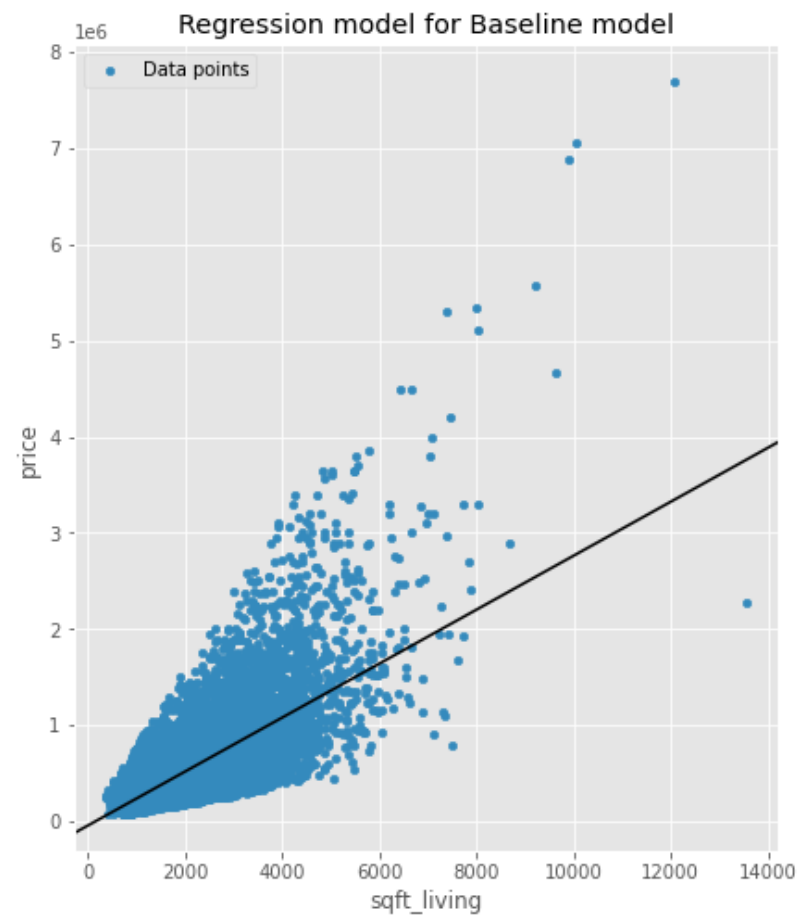
- Compared to grade 3, grades 9,11,12 and 13 have a price increase with grade 13 having the most increase
- Compared to view 0, all views have a price increase with view 4 having the most then view 1

### **Visualizations of the models**

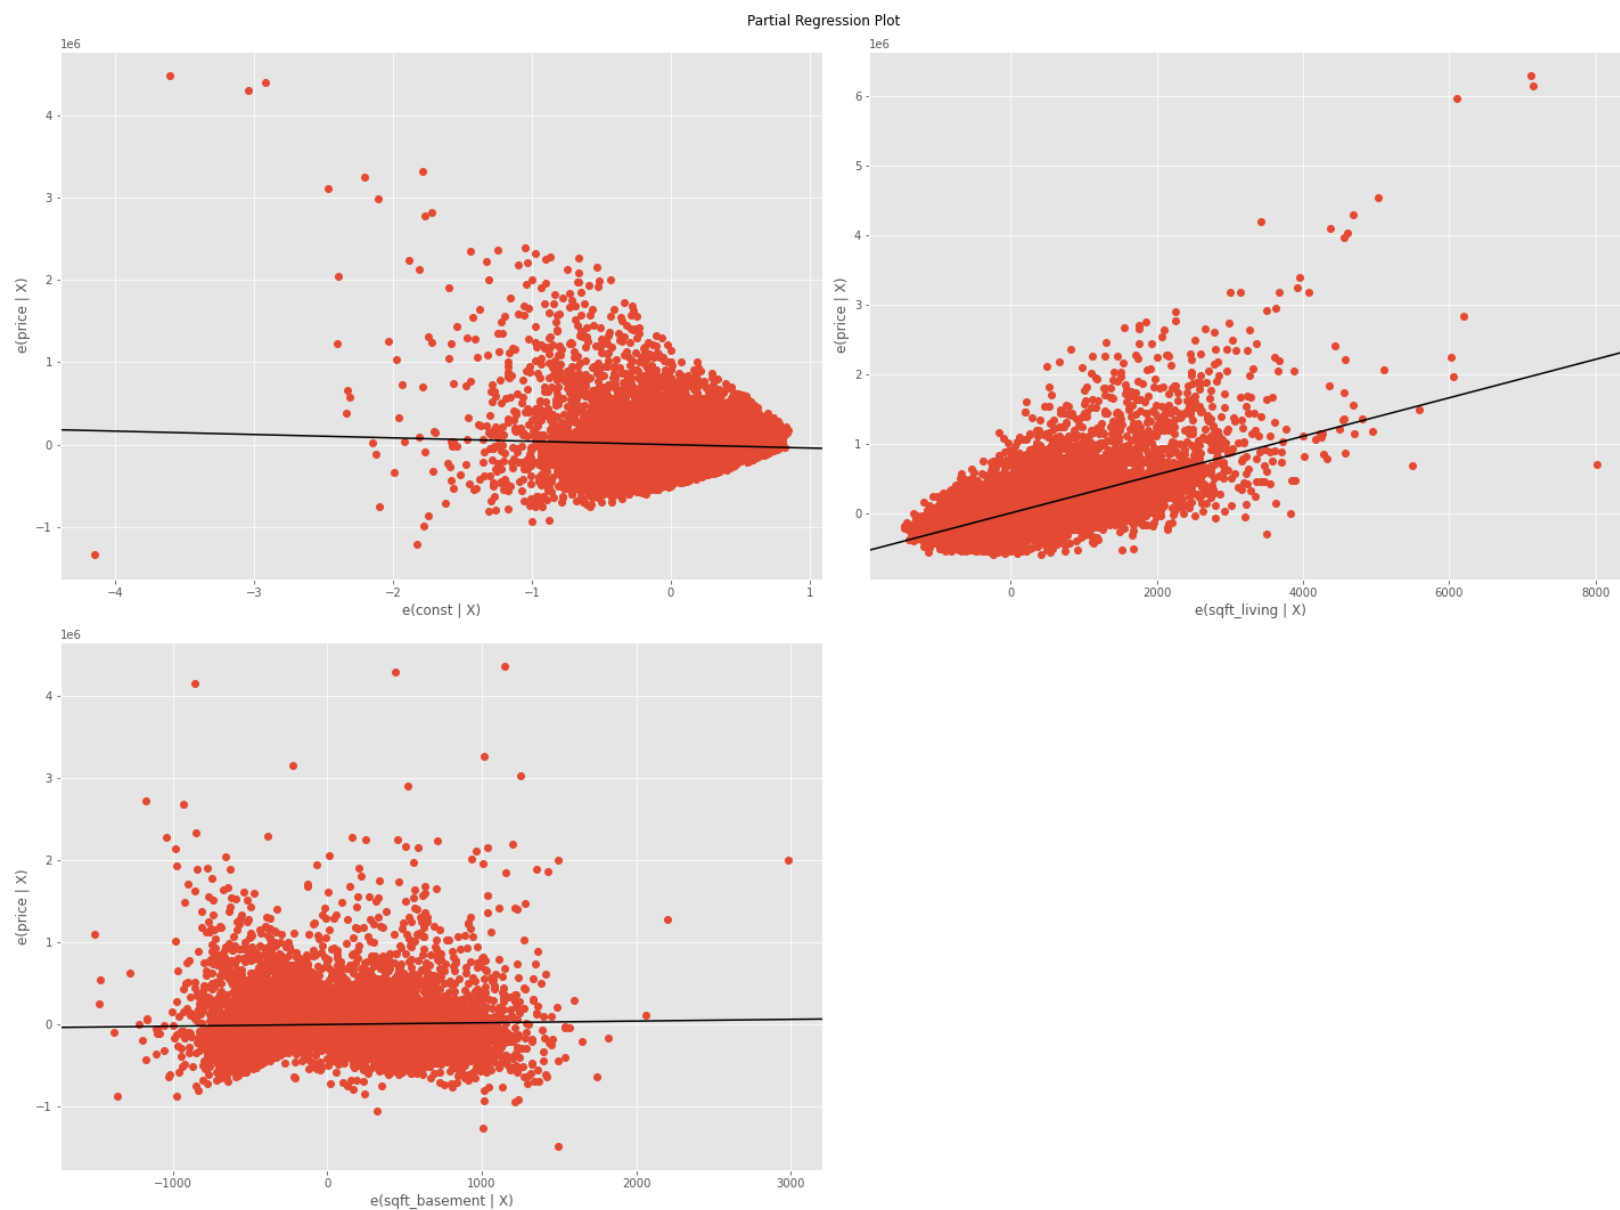


```
In [400]: # Baseline model
# Plot the regression line
fig, ax = plt.subplots(nrows=1,ncols=2,figsize=(15,8))
features_df.plot.scatter(x="sqft_living", y="price", label="Data points", ax=ax[0])
sm.graphics.abline_plot(model_results=model_baseline, label="Regression line", ax=ax[0], color="black")
ax[0].set_title("Regression model for Baseline model")

#Plot the residuals
ax[1].scatter(features_df["price"], model_baseline.resid)
ax[1].axhline(y=0, color="black")
ax[1].set_xlabel("sqft_living")
ax[1].set_ylabel("residuals")
ax[1].set_title("Residuals for baseline model")
plt.savefig("images/Baseline model plots");
```

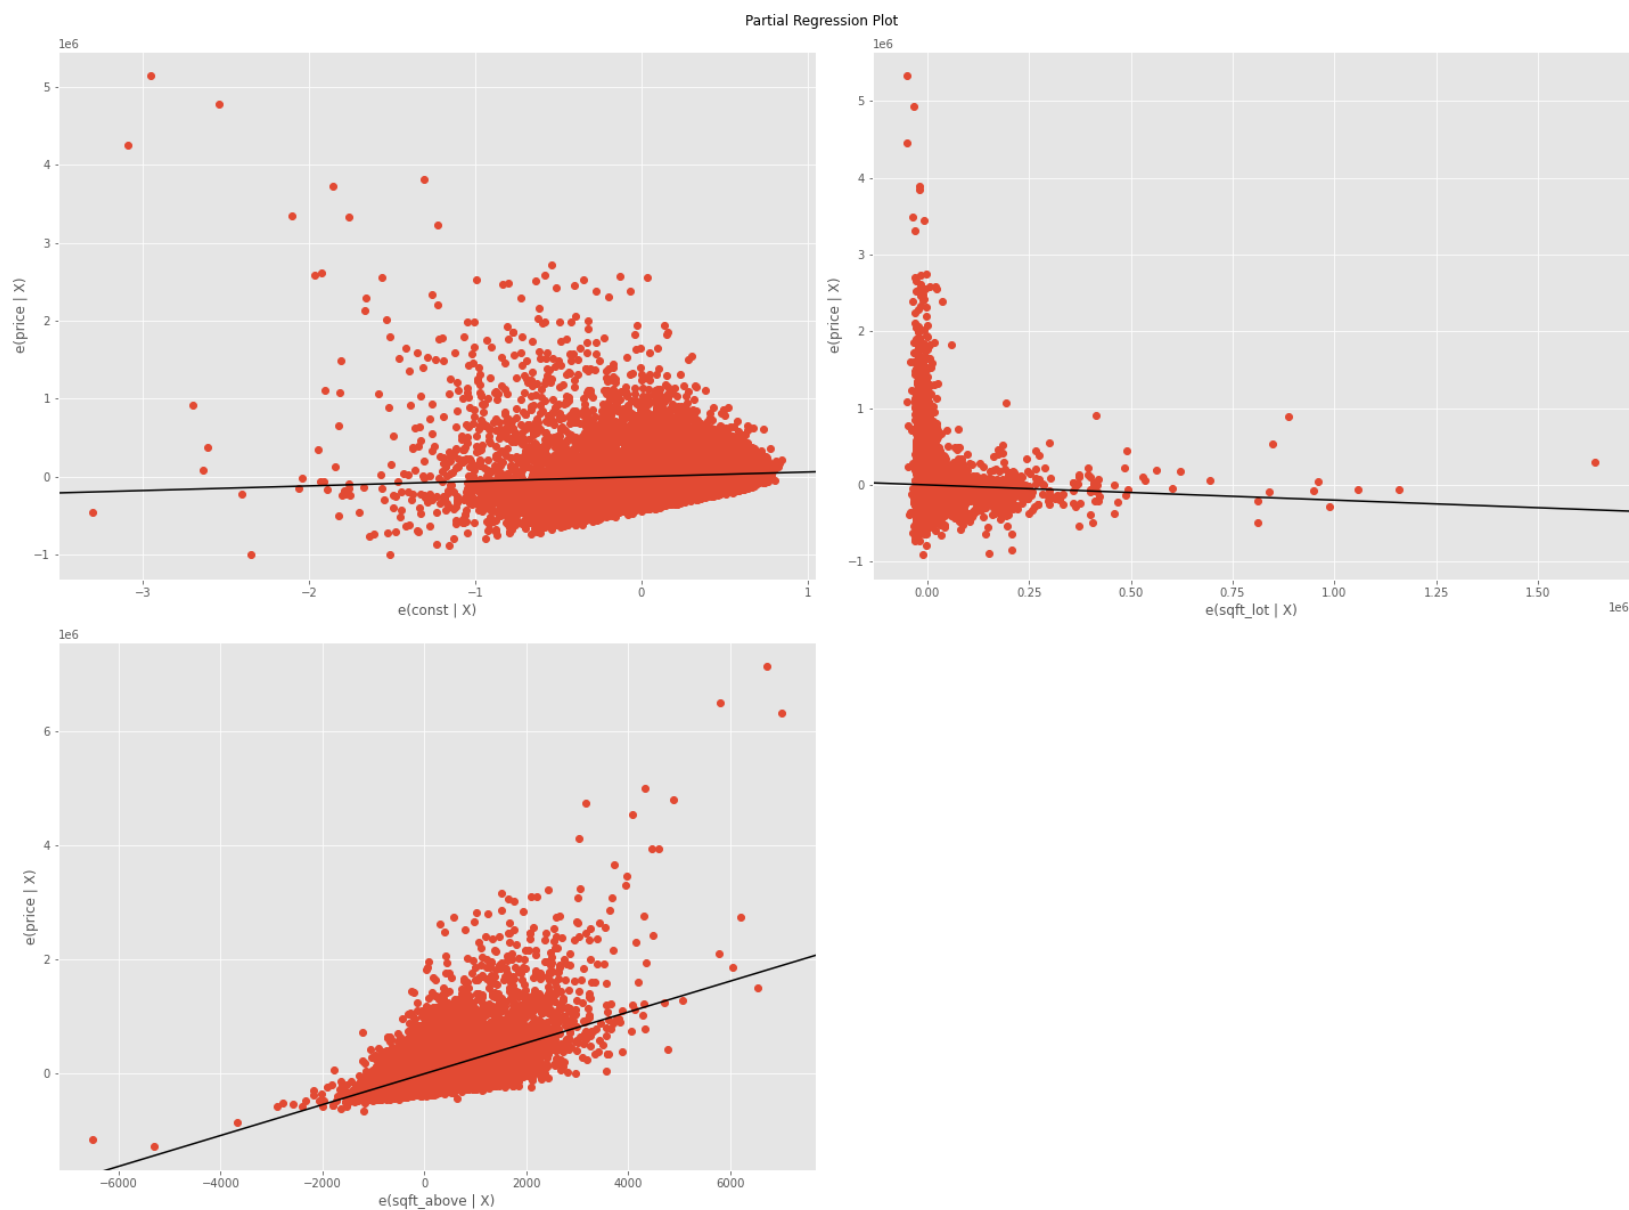


```
In [414]: # Model 1
fig = plt.figure(figsize=(20,15))
sm.graphics.plot_partregress_grid(model1,fig=fig)
plt.savefig("images/model1 partial regression plot");
```



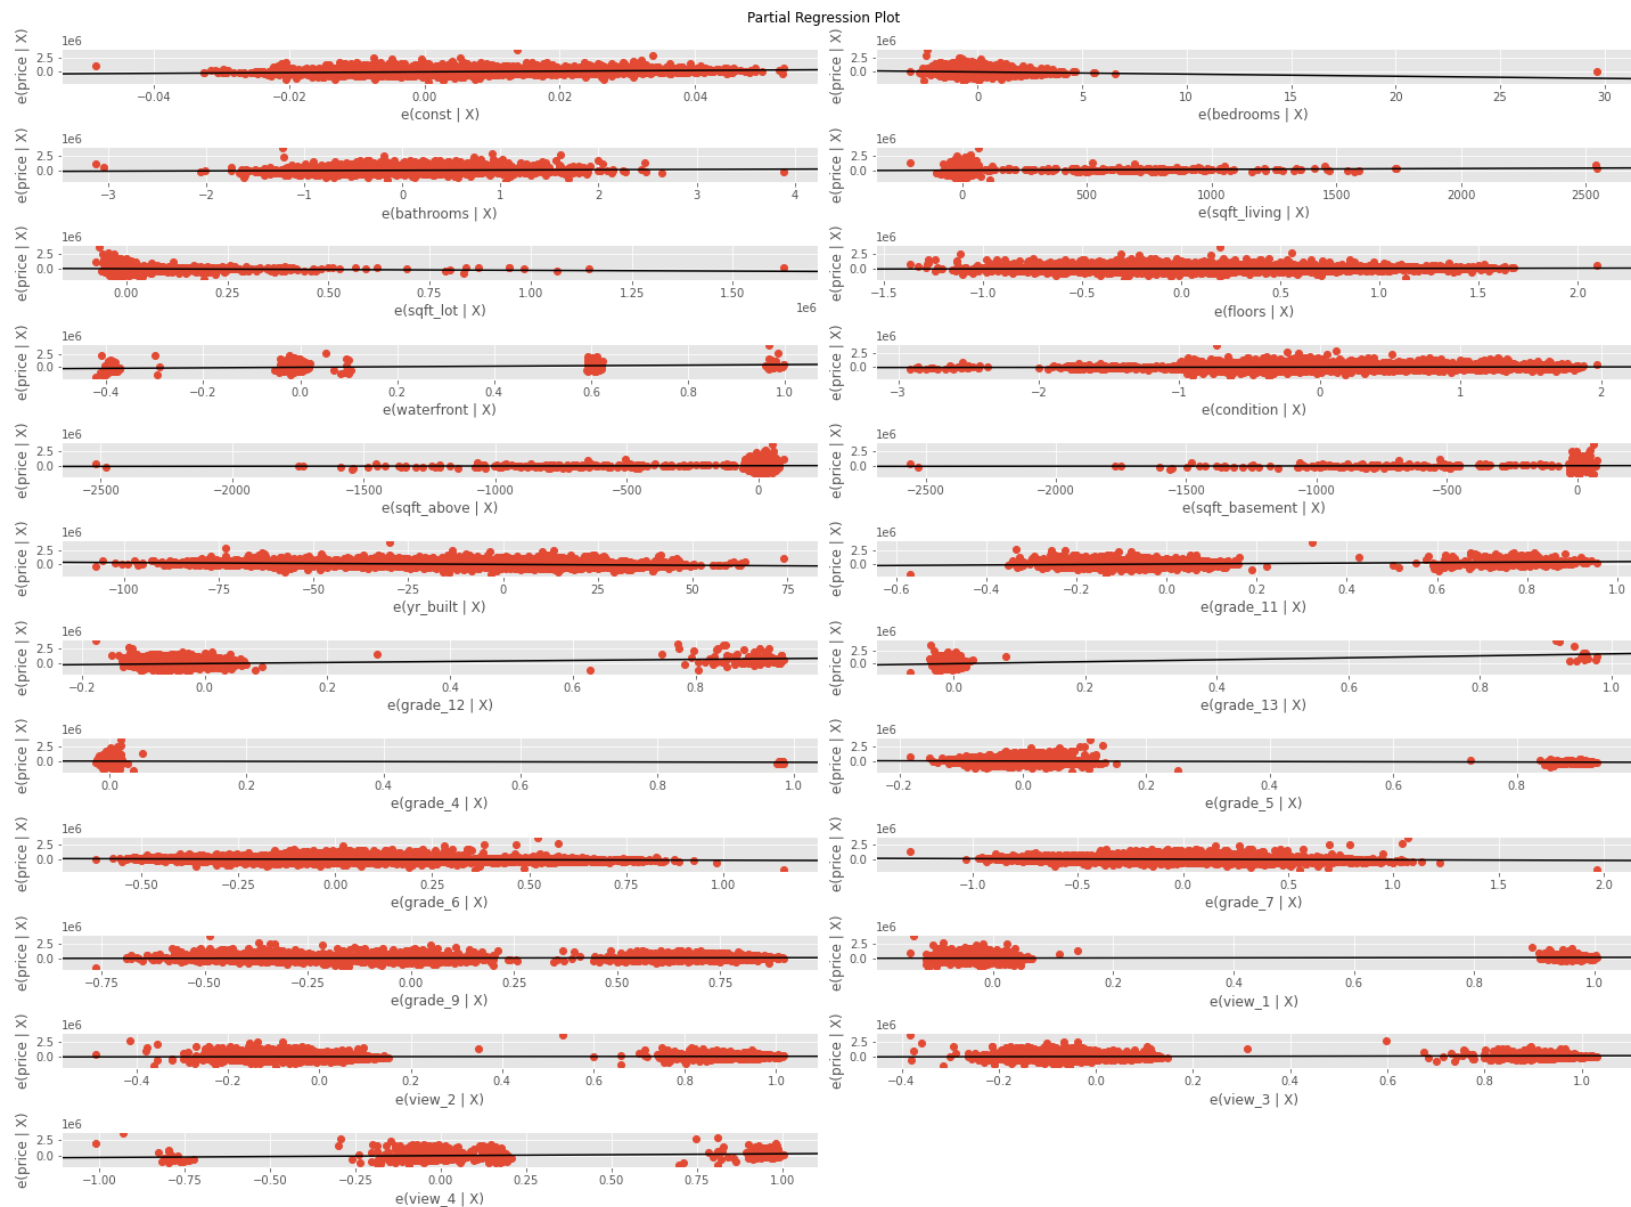


```
In [415]: # Model 2
fig = plt.figure(figsize=(20,15))
sm.graphics.plot_partregress_grid(model2,fig=fig)
plt.savefig("images/model2 partial regression plot")
```





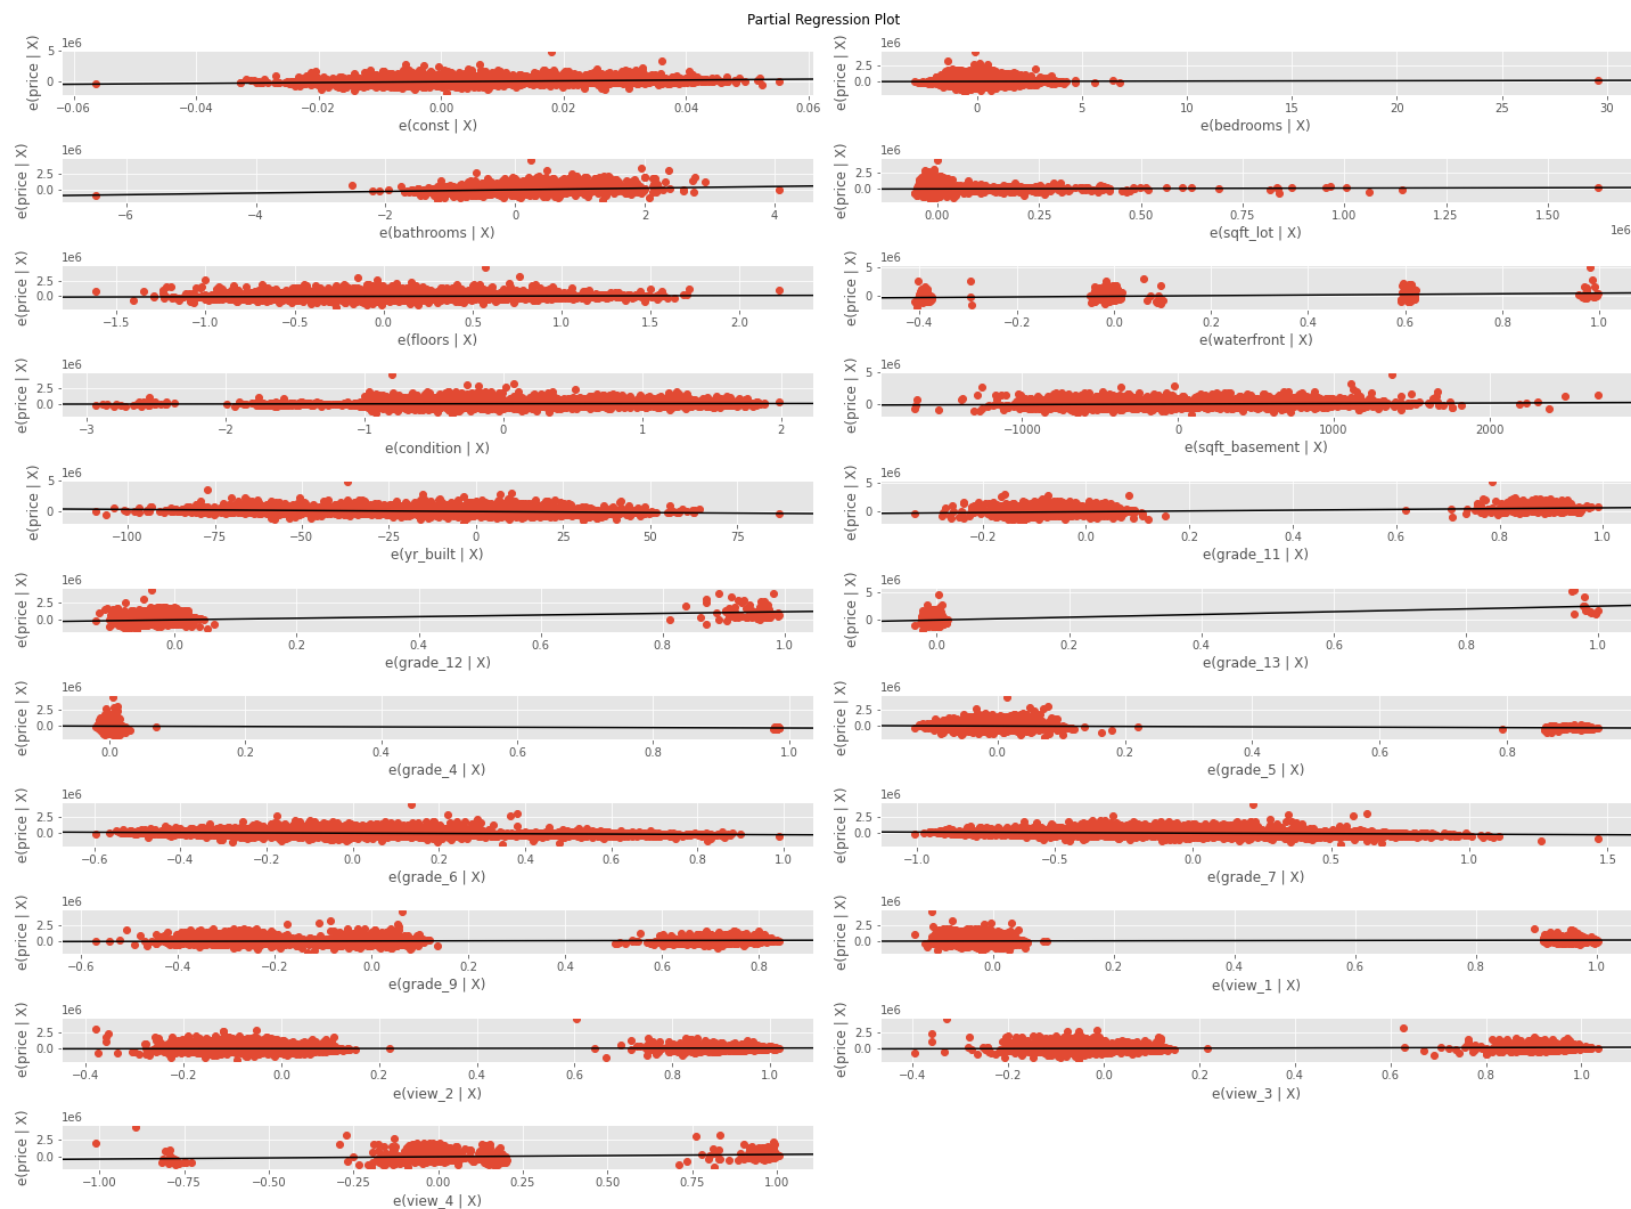
```
In [416]: # Model 3
fig = plt.figure(figsize=(20,15))
sm.graphics.plot_partregress_grid(model3,fig=fig)
plt.savefig("images/model3 partial regression plot");
```







```
In [418]: # Model 4
fig = plt.figure(figsize=(20,15))
sm.graphics.plot_partregress_grid(model4,fig=fig)
plt.savefig("images/model4 partial regression plot");
```

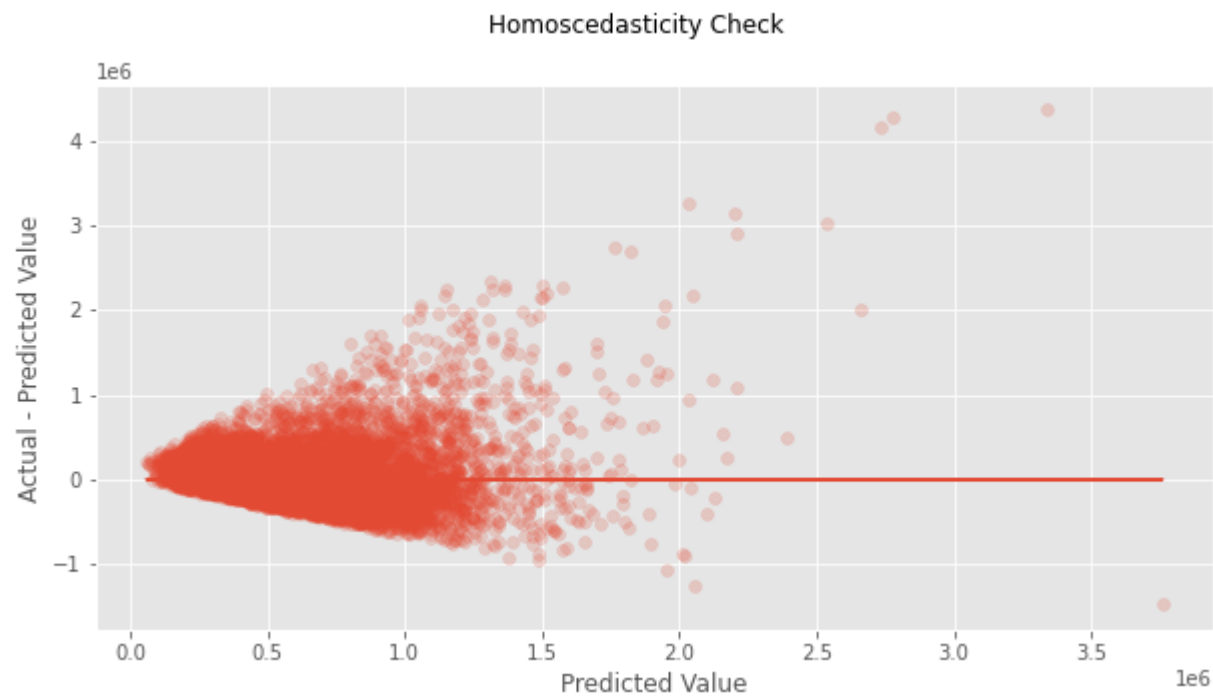


From the above residual plots, it is seen that from the data, the linear regression would be the best fit in comparison to a non-linear one

Homoscedasticity

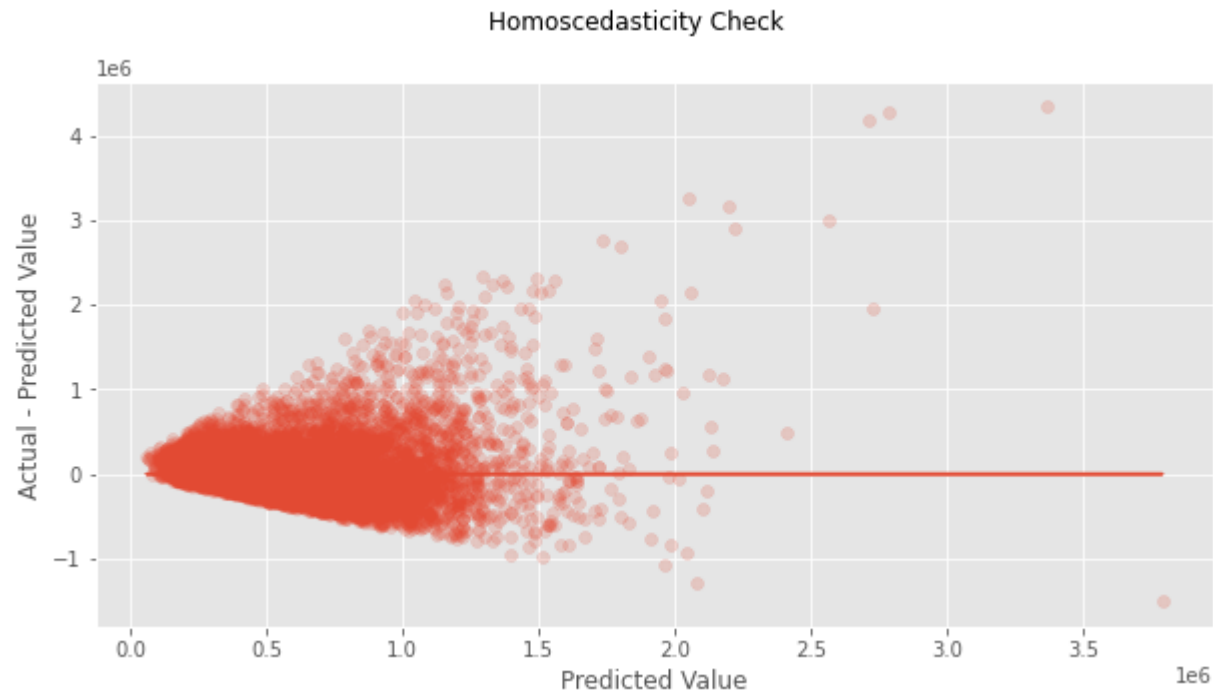
```
In [405]: # Baseline Model
preds = model_baseline.predict(sm.add_constant(X_model_baseline))
residuals = model_baseline.resid

fig, ax = plt.subplots(figsize=(10,5))
fig.suptitle('Homoscedasticity Check')
ax.scatter(preds, residuals, alpha=0.2)
ax.plot(preds, [0 for i in range(len(X_model_baseline))])
ax.set_xlabel("Predicted Value")
ax.set_ylabel("Actual - Predicted Value")
plt.savefig("images/Baseline Model Homoscedasticity check");
```



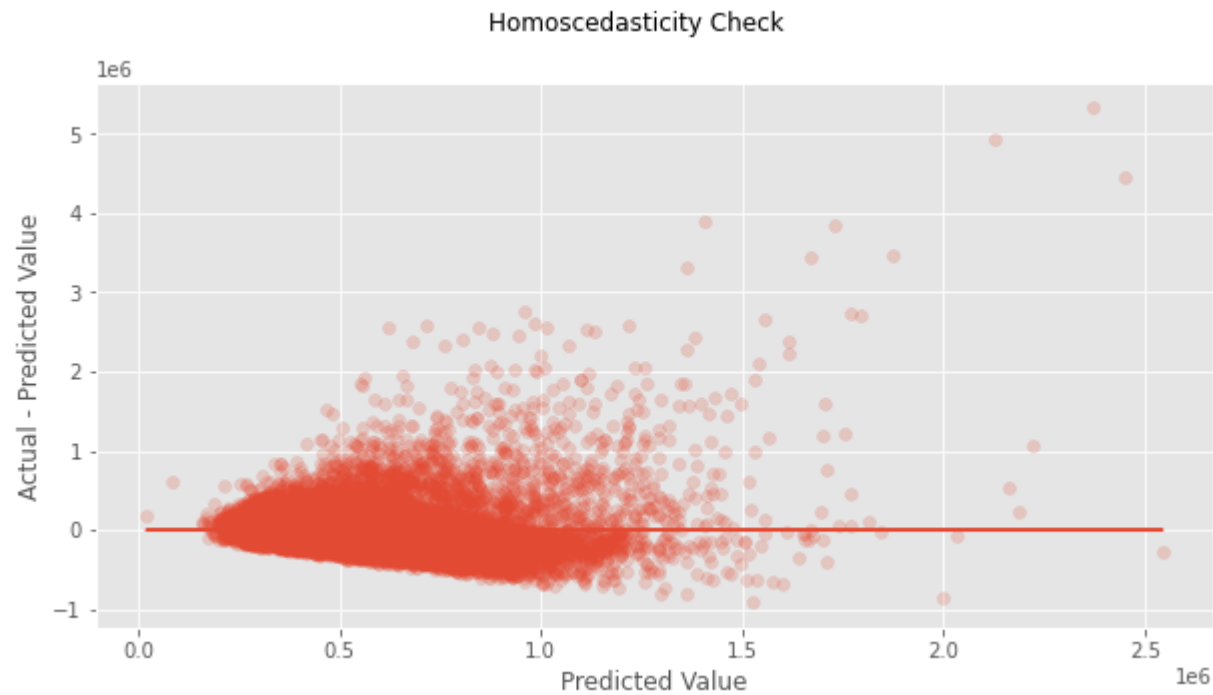
```
In [406]: # Model 1
preds = model1.predict(sm.add_constant(X_model1))
residuals = model1.resid

fig, ax = plt.subplots(figsize=(10,5))
fig.suptitle('Homoscedasticity Check')
ax.scatter(preds, residuals, alpha=0.2)
ax.plot(preds, [0 for i in range(len(X_model1))])
ax.set_xlabel("Predicted Value")
ax.set_ylabel("Actual - Predicted Value")
plt.savefig("images/Model one Homoscedasticity check");
```



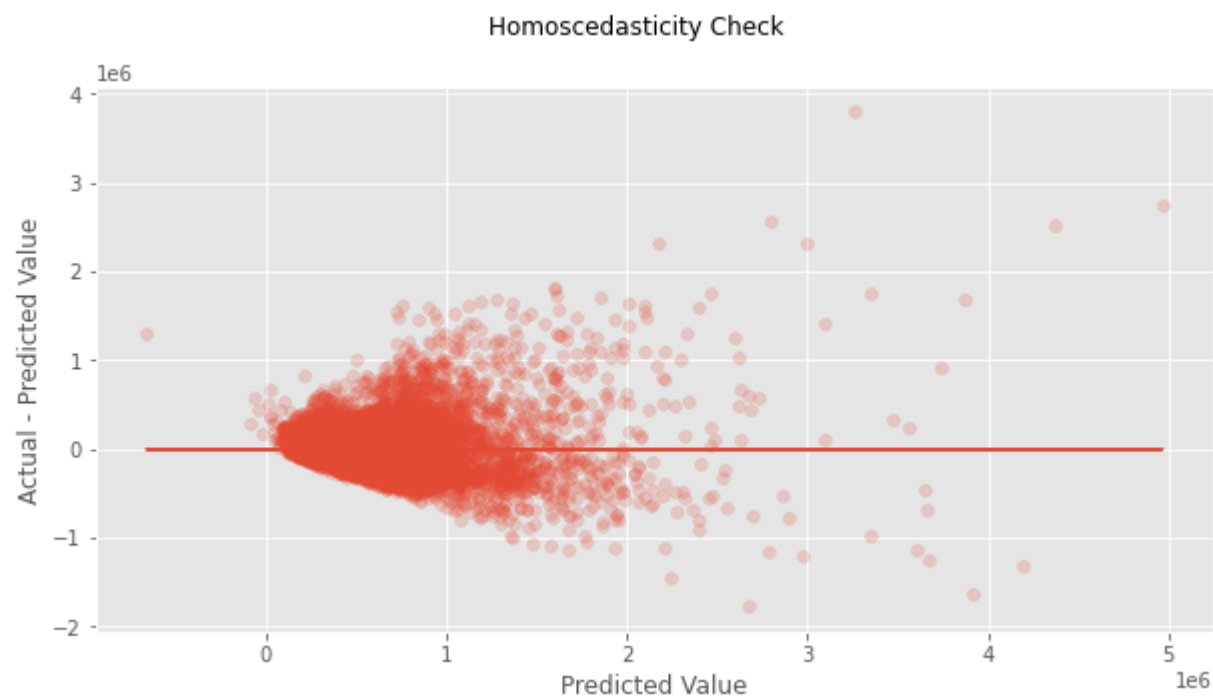
```
In [407]: # Model 2
preds = model2.predict(sm.add_constant(X_model2))
residuals = model2.resid

fig, ax = plt.subplots(figsize=(10,5))
fig.suptitle('Homoscedasticity Check')
ax.scatter(preds, residuals, alpha=0.2)
ax.plot(preds, [0 for i in range(len(X_model2))])
ax.set_xlabel("Predicted Value")
ax.set_ylabel("Actual - Predicted Value")
plt.savefig("images/Model two Homoscedasticity check");
```



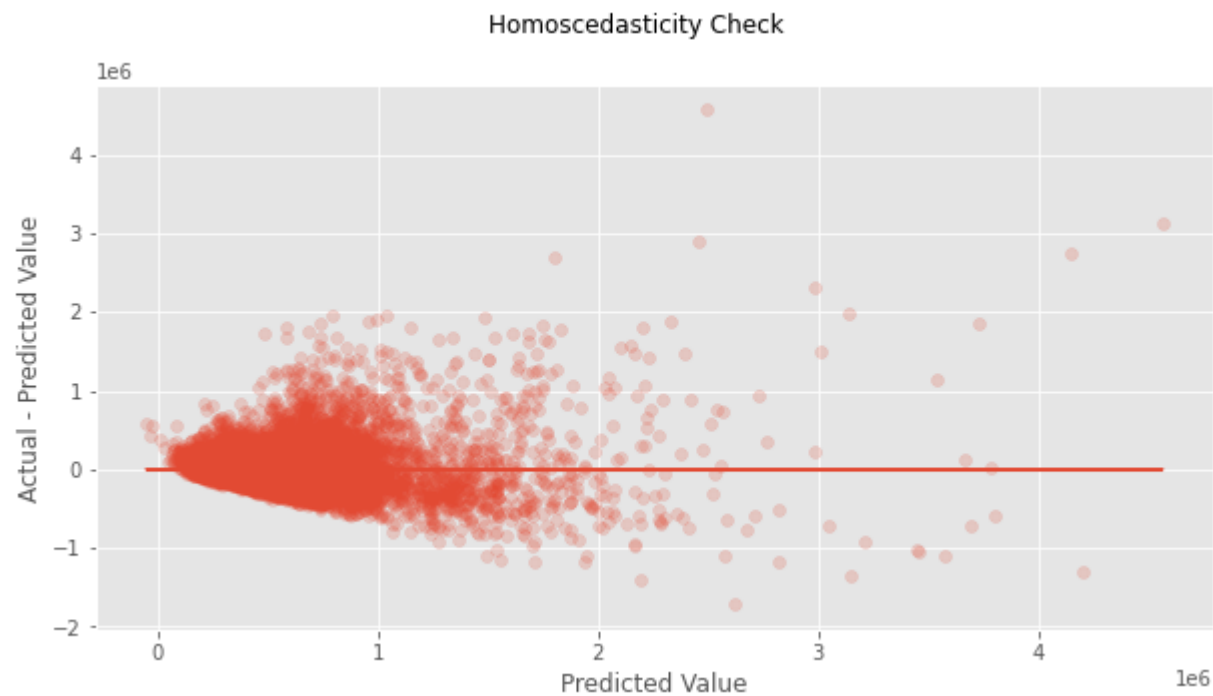
```
In [408]: # Model 3
preds = model3.predict(sm.add_constant(X_model3))
residuals = model3.resid

fig, ax = plt.subplots(figsize=(10,5))
fig.suptitle('Homoscedasticity Check')
ax.scatter(preds, residuals, alpha=0.2)
ax.plot(preds, [0 for i in range(len(X_model3))])
ax.set_xlabel("Predicted Value")
ax.set_ylabel("Actual - Predicted Value")
plt.savefig("images/Model three Homoscedasticity check");
```



```
In [409]: # Model 4
preds = model4.predict(sm.add_constant(X_model4))
residuals = model4.resid

fig, ax = plt.subplots(figsize=(10,5))
fig.suptitle('Homoscedasticity Check')
ax.scatter(preds, residuals, alpha=0.2)
ax.plot(preds, [0 for i in range(len(X_model4))])
ax.set_xlabel("Predicted Value")
ax.set_ylabel("Actual - Predicted Value")
plt.savefig("images/Model four Homoscedasticity check");
```



The models fail the Homoscedasticity check

## Conclusion

### Final model

Model 3 is selected out of all the models as it has a higher r-squared value and a low mean absolute error. The model's characteristics are:

- The overall model explains about 65.4% of the variance in the prices.
- The model is off by about \$141,119 in price
- All coefficients are statistically significant
- Compared to grade 3, grades 4,5,6 and 7 have a price decrease
- Compared to grade 3, grades 9,11,12 and 13 have a price increase with grade 13 having the most increase
- Compared to view 0, views 1,3 and 4 have a price increase with view 4 having the most

## Recommendation

I would recommend the following:

- Houses with a waterfront are set to increase the price by 517,800 \* *The type of view*, compared to having no views would be having a house having a view quality of 1 which is considered fair as it would increase the price by 122,500 and should one want the best view at a quality of 4, they would see an increase in price by \$286,500
- When it comes to building materials, one should go with the rating that is more than average to increase the sale as the highest grade would increase price by 1,923,000



## Limitations and Next steps

The model fits about 65% of the data and additional tests to check for normality and scaling of features should be considered. Also, for future research, deployment of other models besides the regression would be best and the data needs to be analysed while zoning in on specific zipcodes so as to get more accurate and fine-tuned results.



